

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import os
from IPython.display import display, display_html

# Set style & palette
plt.style.use('seaborn-v0_8-whitegrid')
sns.set_theme(style='whitegrid')

# Distinct palette colors
CB_PALETTE = ['#0072B2', '#009E73', '#D55E00', '#CC79A7']
sns.set_palette(CB_PALETTE)

# Custom color mapping for consistency across all plots
COLOR_MAP = {
    'CubeMX v2.2 (Nucleo)': '#0072B2',
    'AI Studio v4.0 (Nucleo)': '#009E73',
    'DevCloud v2.2 (DK)': '#D55E00',
    'DevCloud v4.0 (DK)': '#CC79A7'
}

# Resolved relative to the notebook's location (Report/)
PLOT_DIR = 'plot'
os.makedirs(PLOT_DIR, exist_ok=True)
```

```
In [2]: # Load Baseline files (relative to Report/ directory)
df_base_aistudio = pd.read_csv('../Results/baseline/performance_baseline_res
df_base_cubemx = pd.read_csv('../Results/baseline/performance_baseline_resul
df_base_cloud2 = pd.read_csv('../Results/baseline/performance_baseline_resul
df_base_cloud4 = pd.read_csv('../Results/baseline/performance_baseline_resul

# Load Estimated files (relative to Report/ directory)
df_est_aistudio = pd.read_csv('../Results/estimated/performance_estimated_re
df_est_cubemx = pd.read_csv('../Results/estimated/performance_estimated_resu
df_est_cloud2 = pd.read_csv('../Results/estimated/performance_estimated_resu
df_est_cloud4 = pd.read_csv('../Results/estimated/performance_estimated_resu

# Standardize Model Names across datasets to prevent mismatches
for df in [df_base_aistudio, df_base_cubemx, df_base_cloud2, df_base_cloud4,
           df_est_aistudio, df_est_cubemx, df_est_cloud2, df_est_cloud4]:
    if 'Model Name' in df.columns:
        df['Model Name'] = df['Model Name'].str.replace('conv2d_poll', 'conv

# Data cleaning: convert columns to numeric for baseline data
for df in [df_base_aistudio, df_base_cubemx, df_base_cloud2, df_base_cloud4]:
    if 'Inference Time (ms)' in df.columns:
        df['Inference Time (ms)'] = pd.to_numeric(df['Inference Time (ms)'],
    if 'RAM Usage (KiB)' in df.columns:
        df['RAM Usage (KiB)'] = pd.to_numeric(df['RAM Usage (KiB)'], errors=
    if 'Flash Usage (KiB)' in df.columns:
        df['Flash Usage (KiB)'] = pd.to_numeric(df['Flash Usage (KiB)'], err
```

```
# Map raw model names to formatted labels for presentation
def get_clean_name(model_name):
    if not isinstance(model_name, str):
        return model_name
    name = model_name.replace('baseline_', '').replace('_int8', '')
    if name == 'conv2d_pool':
        return 'Conv2D + MaxPool'
    if name == 'depthwise_3x3':
        return 'Depthwise Conv2D 3x3'
    if name == 'conv2d_3x3':
        return 'Conv2D 32f 3x3'
    if name == 'conv2d_16f_3x3':
        return 'Conv2D 16f 3x3'
    if name == 'conv2d_64f_3x3':
        return 'Conv2D 64f 3x3'
    if name == 'conv2d_1x1':
        return 'Conv2D 32f 1x1'
    if name == 'conv2d_32f_5x5':
        return 'Conv2D 32f 5x5'
    if name == 'conv2d_32f_7x7':
        return 'Conv2D 32f 7x7'
    return name
```

1. Tested Models and Platforms

This section describes the single-layer models selected for benchmarking and the hardware/software configurations tested.

Benchmark Platforms & Configurations:

- **AI Studio v4.0 (Nucleo):** STM32 AI Studio running **AI Core v4.0** compiler on the STM32N6 Nucleo hardware board.
- **CubeMX v2.2 (Nucleo):** STM32CubeMX with X-CUBE-AI running **AI Core v2.2** compiler on the STM32N6 Nucleo hardware board.
- **DevCloud v2.2 (DK):** ST DevCloud running **AI Core v2.2** compiler on the STM32N6 Discovery Kit (DK) board.
- **DevCloud v4.0 (DK):** ST DevCloud running **AI Core v4.0** compiler on the STM32N6 Discovery Kit (DK) board.

```
In [3]: # Compile Table of Tested Models
overview_models = df_base_aistudio[['Model Name', 'Layer Type', 'Filters', 'Clean Name']]
overview_models['Clean Name'] = overview_models['Model Name'].apply(get_clean_name)

# Reorder columns
overview_models = overview_models[['Model Name', 'Clean Name', 'Layer Type', 'Filters']]

# Style and display
styled_models = overview_models.style.set_properties(**{
    'text-align': 'center',
    'font-size': '11pt',
    'border': '1px solid lightgrey'
})
```

```
}).set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#34495e'), ('color',
)]).set_caption("Summary of Tested Models and Operators")

display(styled_models)
```

Summary of Tested Models and Operators

	Model Name	Clean Name	Layer Type	Filters	Kernel Size	IS
0	baseline_conv2d_16f_3x3_int8	Conv2D 16f 3x3	Conv2D	16	3x3	64
1	baseline_conv2d_1x1_int8	Conv2D 32f 1x1	Conv2D (Pointwise)	32	1x1	64
2	baseline_conv2d_32f_5x5_int8	Conv2D 32f 5x5	Conv2D	32	5x5	64
3	baseline_conv2d_32f_7x7_int8	Conv2D 32f 7x7	Conv2D	32	7x7	64
4	baseline_conv2d_3x3_int8	Conv2D 32f 3x3	Conv2D	32	3x3	64
5	baseline_conv2d_64f_3x3_int8	Conv2D 64f 3x3	Conv2D	64	3x3	64
6	baseline_conv2d_pool_int8	Conv2D + MaxPool	Conv2D + MaxPool	16	3x3	64
7	baseline_depthwise_3x3_int8	Depthwise Conv2D 3x3	DepthwiseConv2D	3	3x3	64

2. Inference Time Comparison (Table & Bar Plot)

This section evaluates the execution speed across all four compilers and platforms.

- **FAIL** denotes hardware timeout, out of memory, or model compilation failure.
- The plots use highly distinct, custom colors for visual clarity.

```
In [4]: # Consolidated Inference Time Table
dfs_base = {
    'CubeMX v2.2 (Nucleo)': df_base_cubemx,
    'AI Studio v4.0 (Nucleo)': df_base_aistudio,
    'DevCloud v2.2 (DK)': df_base_cloud2,
    'DevCloud v4.0 (DK)': df_base_cloud4
}

inference_data = []
all_models = df_base_aistudio['Model Name'].unique()

for m in all_models:
```

```

row_data = {'Model': get_clean_name(m)}
for tool_name, df in dfs_base.items():
    row_match = df[df['Model Name'] == m]
    if not row_match.empty:
        val = row_match.iloc[0]['Inference Time (ms)']
        row_data[tool_name] = val
    else:
        row_data[tool_name] = np.nan
inference_data.append(row_data)

df_inference_table = pd.DataFrame(inference_data).set_index('Model')

# Reorder columns
col_order_tools = ['CubeMX v2.2 (Nucleo)', 'AI Studio v4.0 (Nucleo)', 'DevCL
df_inference_table = df_inference_table.reindex(columns=col_order_tools)

# Sort index logically by model size
model_order = [
    'Depthwise Conv2D 3x3', 'Conv2D 16f 3x3', 'Conv2D + MaxPool',
    'Conv2D 32f 1x1', 'Conv2D 32f 3x3', 'Conv2D 32f 5x5',
    'Conv2D 64f 3x3', 'Conv2D 32f 7x7'
]
df_inference_table = df_inference_table.reindex([m for m in model_order if m

# Formatted table for printing
df_inference_print = df_inference_table.apply(lambda col: col.map(lambda x:

styled_inference = df_inference_print.style.set_properties(**{
    'text-align': 'center',
    'font-size': '11pt',
    'border': '1px solid lightgrey'
}).set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#2980b9'), ('color',
)].set_caption("Measured Inference Time (ms) across Compilers and Boards")

display(styled_inference)

# Plotting the inference times as a bar plot
plt.figure(figsize=(14, 7))
data_list = []
for tool, df in dfs_base.items():
    for _, row in df.iterrows():
        clean_model = get_clean_name(row['Model Name'])
        val = row['Inference Time (ms)']
        data_list.append({
            'Tool': tool,
            'Model Name': clean_model,
            'Inference Time (ms)': val
        })
df_all = pd.DataFrame(data_list)


sns.barplot(data=df_all, x='Model Name', y='Inference Time (ms)', hue='Tool')
plt.title('Inference Time (ms) per Layer across Tools & Platforms', fontweig
plt.ylabel('Inference Time (ms) [Linear Scale]', fontsize=12)
plt.xlabel('Layer Name', fontsize=12)
plt.xticks(rotation=30, ha='right')

```

```
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.legend(title='Platform & Tool', frameon=True)
plt.tight_layout()
plt.savefig(f'{PLOT_DIR}/11_inference_time_bar_plot.png', dpi=300)
plt.show()
```

Measured Inference Time (ms) across Compilers and Boards

	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Model				
Depthwise Conv2D 3x3	0.042 ms	0.053 ms	0.050 ms	0.031 ms
Conv2D 16f 3x3	0.079 ms	0.122 ms	0.087 ms	0.086 ms
Conv2D + MaxPool	0.080 ms	0.121 ms	0.087 ms	0.088 ms
Conv2D 32f 1x1	0.146 ms	0.373 ms	0.153 ms	0.288 ms
Conv2D 32f 3x3	0.147 ms	0.208 ms	0.155 ms	0.155 ms
Conv2D 32f 5x5	0.154 ms	0.219 ms	0.161 ms	0.161 ms
Conv2D 64f 3x3	0.284 ms	0.383 ms	0.291 ms	0.291 ms
Conv2D 32f 7x7	0.462 ms	FAIL	0.473 ms	FAIL



3. Memory Footprint: Compiler Estimated vs. Actual Hardware RAM Usage

Here we compare the **Estimated RAM** against the **Actual RAM** using distinct color scales:

- **Logarithmic scale** is used to handle large variations (e.g. from 12 KiB to 1548 KiB).
- **Shared Y-axis scale** is applied to directly expose discrepancies, highlighting cases where compilers severely underestimate actual memory limits (e.g. 7x7 layer on CubeMX).

```
In [5]: # Prepare RAM data
ram_data_est = []
ram_data_act = []

dfs_est = {
    'CubeMX v2.2 (Nucleo)': df_est_cubemx,
    'AI Studio v4.0 (Nucleo)': df_est_aistudio,
    'DevCloud v2.2 (DK)': df_est_cloud2,
    'DevCloud v4.0 (DK)': df_est_cloud4
}

for m in all_models:
    clean_m = get_clean_name(m)

    # Estimated
    est_row = {'Model': clean_m}
    for tool_name, df in dfs_est.items():
        row_match = df[df['Model Name'] == m]
        val = row_match.iloc[0]['Estimated RAM Usage (KiB)'] if not row_match.empty else 0
        est_row[tool_name] = val
    ram_data_est.append(est_row)

    # Actual
    act_row = {'Model': clean_m}
    for tool_name, df in dfs_base.items():
        row_match = df[df['Model Name'] == m]
        val = row_match.iloc[0]['RAM Usage (KiB)'] if not row_match.empty else 0
        act_row[tool_name] = val
    ram_data_act.append(act_row)

df_ram_est = pd.DataFrame(ram_data_est).set_index('Model').reindex(columns=columns)
df_ram_act = pd.DataFrame(ram_data_act).set_index('Model').reindex(columns=columns)

df_ram_est = df_ram_est.reindex([m for m in model_order if m in df_ram_est.index])
df_ram_act = df_ram_act.reindex([m for m in model_order if m in df_ram_act.index])

# Plotting side-by-side with custom colors
fig, axes = plt.subplots(1, 2, figsize=(18, 7), sharey=True)
colors_ram = [COLOR_MAP[col] for col in df_ram_est.columns]

df_ram_est.plot(kind='bar', ax=axes[0], width=0.8, color=colors_ram)
axes[0].set_title('Estimated RAM Usage (KiB) - Compiler Prediction', fontweight='bold')
axes[0].set_ylabel('RAM (KiB) [Log Scale]', fontsize=11)
axes[0].set_yscale('log')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=30, ha='right')
axes[0].grid(True, which="both", ls="--", alpha=0.5)

df_ram_act.plot(kind='bar', ax=axes[1], width=0.8, color=colors_ram)
axes[1].set_title('Actual RAM Usage (KiB) - Measured on Hardware', fontweight='bold')
```

```

axes[1].set_yscale('log')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=30, ha='right')
axes[1].grid(True, which="both", ls="--", alpha=0.5)

# Set same y-limit
max_val = max(df_ram_est.max().max(), df_ram_act.max().max())
axes[0].set_ylim(1, max_val * 1.5)
axes[1].set_ylim(1, max_val * 1.5)


plt.suptitle('RAM Footprint: Compiler Estimated vs. Actual Hardware Usage (Log Scale)')
plt.tight_layout()
plt.savefig(f'{PLOT_DIR}/9_ram_estimated_vs_actual_side_by_side.png', dpi=300)
plt.show()

# Display tables stacked vertically for proper PDF formatting (avoid horizontal scrolling)
html_est = df_ram_est.style.set_caption("Estimated RAM Usage (KiB)").format(
    {'selector': 'th', 'props': [('background-color', '#8e44ad'), ('color', 'white')]}
).to_html()

html_act = df_ram_act.style.set_caption("Actual RAM Usage (KiB)").format(
    {'selector': 'th', 'props': [('background-color', '#2c3e50'), ('color', 'white')]}
).to_html()

display_html(html_est, raw=True)
display_html("<br>", raw=True)
display_html(html_act, raw=True)

```



Estimated RAM Usage (KiB)				
	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Model				
Depthwise Conv2D 3x3	24.00	12.00	24.00	12.00
Conv2D 16f 3x3	12.02	12.02	12.02	12.02
Conv2D + MaxPool	12.02	12.02	12.02	12.02
Conv2D 32f 1x1	12.03	140.00	12.03	140.00
Conv2D 32f 3x3	12.03	12.03	12.03	12.03
Conv2D 32f 5x5	12.03	12.03	12.03	12.03
Conv2D 64f 3x3	12.06	12.06	12.06	12.06
Conv2D 32f 7x7	1292.00	1546.24	1290.24	1546.24

Actual RAM Usage (KiB)				
	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Model				
Depthwise Conv2D 3x3	24.00	17.34	13.07	12.00
Conv2D 16f 3x3	12.02	76.28	77.20	12.02
Conv2D + MaxPool	12.02	19.30	20.22	12.02
Conv2D 32f 1x1	12.03	130.26	130.34	140.00
Conv2D 32f 3x3	12.03	140.42	142.28	12.03
Conv2D 32f 5x5	12.03	140.61	145.47	12.03
Conv2D 64f 3x3	12.06	268.70	272.44	12.06
Conv2D 32f 7x7	1292.00	140.92	146.06	1548.00

4. Operation Complexity: Compiler Estimated vs. Actual Hardware MACC

Below is the comparison of **Estimated MACC** vs. **Actual MACC** count for all models across compilers. Linear scale is used for readability.

In [6]:

```
# Prepare MACC data
macc_data_est = []
```



```

macc_data_act = []

for m in all_models:
    clean_m = get_clean_name(m)

    # Estimated
    est_row = {'Model': clean_m}
    for tool_name, df in dfs_est.items():
        row_match = df[df['Model Name'] == m]
        val = row_match.iloc[0]['Estimated MACC'] if not row_match.empty else
        est_row[tool_name] = pd.to_numeric(val, errors='coerce')
    macc_data_est.append(est_row)

    # Actual
    act_row = {'Model': clean_m}
    for tool_name, df in dfs_base.items():
        row_match = df[df['Model Name'] == m]
        val = row_match.iloc[0]['MACC'] if not row_match.empty else np.nan
        act_row[tool_name] = pd.to_numeric(val, errors='coerce')
    macc_data_act.append(act_row)

df_macc_est = pd.DataFrame(macc_data_est).set_index('Model').reindex(columns
df_macc_act = pd.DataFrame(macc_data_act).set_index('Model').reindex(columns

df_macc_est = df_macc_est.reindex([m for m in model_order if m in df_macc_es
df_macc_act = df_macc_act.reindex([m for m in model_order if m in df_macc_ac

# Plotting MACC side-by-side with custom colors
fig, axes = plt.subplots(1, 2, figsize=(18, 7), sharey=True)
colors_macc = [COLOR_MAP[col] for col in df_macc_est.columns]

df_macc_est.plot(kind='bar', ax=axes[0], width=0.8, color=colors_macc)
axes[0].set_title('Estimated MACC (Compiler)', fontweight='bold', fontsize=1
axes[0].set_ylabel('MAC Operations [Linear Scale]', fontsize=11)
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=30, ha='right')
axes[0].grid(True, which="both", ls="--", alpha=0.5)

df_macc_act.plot(kind='bar', ax=axes[1], width=0.8, color=colors_macc)
axes[1].set_title('Actual MACC (Hardware)', fontweight='bold', fontsize=12)
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=30, ha='right')
axes[1].grid(True, which="both", ls="--", alpha=0.5)

# Set same y-limit
max_macc = max(df_macc_est.max().max(), df_macc_act.max().max())
axes[0].set_ylim(0, max_macc * 1.1)
axes[1].set_ylim(0, max_macc * 1.1)


plt.suptitle('Operator Complexity: Estimated vs. Actual MACC (Linear Scale)'
plt.tight_layout()
plt.savefig(f'{PLOT_DIR}/10_macc_estimated_vs_actual_side_by_side.png', dpi=
plt.show()

# Display tables stacked vertically for proper PDF formatting
html_macc_est = df_macc_est.style.set_caption("Estimated MACC").format(preci
    {'selector': 'th', 'props': [('background-color', '#16a085'), ('color',
    ]}).to_html()

```

```
html_macc_act = df_macc_act.style.set_caption("Actual MACC").format(precision=2,
    {'selector': 'th', 'props': [('background-color', '#2c3e50'), ('color', 'white')]}).to_html()

display_html(html_macc_est, raw=True)
display_html("<br>", raw=True)
display_html(html_macc_act, raw=True)
```



Estimated MACC

	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Model				
Depthwise Conv2D 3x3	135171	116736	0	337920
Conv2D 16f 3x3	1900560	1802240	0	1802240
Conv2D + MaxPool	1916944	1802240	0	1802240
Conv2D 32f 1x1	655392	458752	0	458752
Conv2D 32f 3x3	3801120	3604480	0	3604480
Conv2D 32f 5x5	10092576	9961472	0	9961472
Conv2D 64f 3x3	7602240	7208960	0	7208960
Conv2D 32f 7x7	19529760	19660800	0	19660800

Actual MACC

Model	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Depthwise Conv2D 3x3	135171	122883	122883	337920
Conv2D 16f 3x3	1900560	1835024	1835024	1802240
Conv2D + MaxPool	1916944	1851408	1851408	1802240
Conv2D 32f 1x1	655392	524320	524320	458752
Conv2D 32f 3x3	3801120	3670048	3670048	3604480
Conv2D 32f 5x5	10092576	9961504	9961504	9961472
Conv2D 64f 3x3	7602240	7340096	7340096	7208960
Conv2D 32f 7x7	19529760	19398688	19398688	19660800

5. Hardware Epochs Partitioning

On the STM32N6 NPU, models can be split into multiple "hardware epochs" to fit within the internal SRAM. A higher number of epochs indicates that intermediate activations or weights had to be tiled and swapped back and forth. The table below lists the epoch partitioning strategy for each tool/compiler.

```
In [7]: # Extract hardware epochs from notes
epochs_data = []

def parse_epochs(notes):
    if pd.isna(notes):
        return 'N/A'
    notes_lower = str(notes).lower()
    if 'fail' in notes_lower or 'timeout' in notes_lower:
        return 'FAIL'

    if '1 single' in notes_lower or '1 hardware' in notes_lower or 'single h
        return '1'
    if '2' in notes_lower:
        return '2'
    if '3' in notes_lower:
        return '3'
    if '4' in notes_lower:
        return '4'

    import re
    nums = re.findall(r'\d+', notes)
    if nums:
        return nums[0]
```

```
    return notes

for tool, df in dfs_base.items():
    for _, row in df.iterrows():
        clean_model = get_clean_name(row['Model Name'])
        epoch_str = parse_epochs(row['Notes'])
        epochs_data.append({
            'Tool': tool,
            'Model Name': clean_model,
            'Hardware Epochs': epoch_str
        })

df_epochs = pd.DataFrame(epochs_data)
df_epochs_pivot = df_epochs.pivot(index='Model Name', columns='Tool', values='Hardware Epochs')
df_epochs_pivot = df_epochs_pivot.reindex(columns=col_order_tools)
df_epochs_pivot = df_epochs_pivot.reindex(index=[c for c in model_order if c in df_epochs_pivot.index])

styled_epochs = df_epochs_pivot.style.set_properties(**{
    'text-align': 'center',
    'font-size': '11pt',
    'border': '1px solid lightgrey'
}).set_table_styles([
    {'selector': 'th', 'props': [('background-color', '#d35400'), ('color', 'white')]}
]).set_caption("Hardware Epochs Partitioning by Layer & Tool")

display(styled_epochs)
```

Hardware Epochs Partitioning by Layer & Tool

Tool	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Model Name				
Depthwise Conv2D 3x3	2	1	2	1
Conv2D 16f 3x3	1	1	1	1
Conv2D + MaxPool	1	1	1	1
Conv2D 32f 1x1	1	2	1	2
Conv2D 32f 3x3	1	1	1	1
Conv2D 32f 5x5	1	1	1	1
Conv2D 64f 3x3	1	1	1	1
Conv2D 32f 7x7	4	FAIL	4	FAIL

6. Comprehensive Platform Statistics & Distributions

This section provides a detailed statistical comparison to identify the overall fastest and most memory-efficient compiler and hardware platform.

- **Summary Tables:** Tables computing Count, Mean, Std Dev, Variance, Min, Medians, Percentiles (25%, 75%, 90%), Max, and Success Rate.
- **Side-by-Side Distribution Box Plots:** Displays the actual variability, standard deviation, and any performance outliers for both Latency (linear scale) and RAM Memory allocation (logarithmic scale) with custom palettes and individual data points overlaid.

```
In [8]: # Compute detailed statistics for successful runs
stats_time_list = []
stats_ram_list = []

for tool, df in dfs_base.items():
    times = df['Inference Time (ms)'].dropna()
    rams = df['RAM Usage (KiB)'].dropna()

    # Calculate Latency Stats
    t_stats = times.describe()
    t_stats['variance'] = times.var()
    t_stats['90%'] = times.quantile(0.9)
    t_stats['success_rate (%)'] = (df['Inference Time (ms)'].notna().sum() /
    stats_time_list.append(pd.DataFrame(t_stats).rename(columns={'Inference

    # Calculate RAM Stats
    r_stats = rams.describe()
    r_stats['variance'] = rams.var()
    r_stats['90%'] = rams.quantile(0.9)
    stats_ram_list.append(pd.DataFrame(r_stats).rename(columns={'RAM Usage (

df_stats_time = pd.concat(stats_time_list, axis=1).reindex(columns=col_order)
df_stats_ram = pd.concat(stats_ram_list, axis=1).reindex(columns=col_order_t

stats_rename = {
    'count': 'Successful Runs (count)',
    'mean': 'Mean / Media (avg)',
    'std': 'Std Dev / Deviaz. Std',
    'variance': 'Variance / Varianza',
    'min': 'Minimum / Minimo',
    '25%': '25th Percentile (P25)',
    '50%': 'Median / Mediana (P50)',
    '75%': '75th Percentile (P75)',
    '90%': '90th Percentile (P90)',
    'max': 'Maximum / Massimo',
    'success_rate (%)': 'Success Rate (%)'
}

df_stats_time.rename(index=stats_rename, inplace=True)
df_stats_ram.rename(index=stats_rename, inplace=True)

# Styled HTML tables displayed stacked vertically for proper PDF width compl
html_stats_time = df_stats_time.style.set_caption("Latency (Inference Time)
    {'selector': 'th', 'props': [('background-color', '#2c3e50'), ('color',
)].to_html()
```

```

html_stats_ram = df_stats_ram.style.set_caption("RAM Memory Allocation Summa
    {'selector': 'th', 'props': [('background-color', '#2c3e50'), ('color',
    ]).to_html()

display_html(html_stats_time, raw=True)
display_html("<br>", raw=True)
display_html(html_stats_ram, raw=True)

# Plotting side-by-side boxplots for distributions
fig, axes = plt.subplots(1, 2, figsize=(18, 7))

# Latency Boxplot (linear scale)
sns.boxplot(
    data=df_all.dropna(subset=['Inference Time (ms)']),
    x='Tool',
    y='Inference Time (ms)',
    ax=axes[0],
    palette=COLOR_MAP,
    order=col_order_tools
)
sns.stripplot(
    data=df_all.dropna(subset=['Inference Time (ms)']),
    x='Tool',
    y='Inference Time (ms)',
    ax=axes[0],
    color='black',
    alpha=0.6,
    size=6,
    jitter=0.15,
    order=col_order_tools
)
axes[0].set_title('Inference Time Distribution (ms) - Lower is Faster', font
axes[0].set_ylabel('Inference Time (ms)', fontsize=11)
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=15, ha='right')
axes[0].grid(axis='y', linestyle='--', alpha=0.7)

# Prepare RAM list for plotting
ram_plot_data = []
for tool, df in dfs_base.items():
    for _, row in df.iterrows():
        ram_plot_data.append({
            'Tool': tool,
            'RAM Usage (KiB)': row['RAM Usage (KiB)']
        })
df_ram_all = pd.DataFrame(ram_plot_data)

# RAM Boxplot (log scale for high range visibility)
sns.boxplot(
    data=df_ram_all.dropna(subset=['RAM Usage (KiB)']),
    x='Tool',
    y='RAM Usage (KiB)',
    ax=axes[1],
    palette=COLOR_MAP,
    order=col_order_tools
)
sns.stripplot(

```

```
data=df_ram_all.dropna(subset=['RAM Usage (KiB)']),
x='Tool',
y='RAM Usage (KiB)',
ax=axes[1],
color='black',
alpha=0.6,
size=6,
jitter=0.15,
order=col_order_tools
)
axes[1].set_title('RAM Allocation Distribution (KiB) - Lower is More Efficient')
axes[1].set_ylabel('RAM (KiB) [Log Scale]', fontsize=11)
axes[1].set_yscale('log')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=15, ha='right')
axes[1].grid(axis='y', linestyle='--', alpha=0.7)

plt.suptitle('Platform Benchmark Distributions: Latency vs. Memory Footprint')
plt.tight_layout()
plt.savefig(f'{PLOT_DIR}/12_latency_vs_ram_distributions.png', dpi=300)
plt.show()
```

Latency (Inference Time) Summary Statistics

	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Successful Runs (count)	8.0000	7.0000	8.0000	7.0000
Mean / Media (avg)	0.1742	0.2113	0.1821	0.1571
Std Dev / Deviaz. Std	0.1374	0.1271	0.1384	0.1007
Minimum / Minimo	0.0420	0.0530	0.0500	0.0310
25th Percentile (P25)	0.0798	0.1215	0.0870	0.0870
Median / Mediana (P50)	0.1465	0.2080	0.1540	0.1550
75th Percentile (P75)	0.1865	0.2960	0.1935	0.2245
Maximum / Massimo	0.4620	0.3830	0.4730	0.2910
Variance / Varianza	0.0189	0.0161	0.0191	0.0101
90th Percentile (P90)	0.3374	0.3770	0.3456	0.2892
Success Rate (%)	100.0000	87.5000	100.0000	87.5000

RAM Memory Allocation Summary Statistics

	CubeMX v2.2 (Nucleo)	AI Studio v4.0 (Nucleo)	DevCloud v2.2 (DK)	DevCloud v4.0 (DK)
Successful Runs (count)	8.00	8.00	8.00	8.00
Mean / Media (avg)	173.52	116.73	118.39	220.02
Std Dev / Deviaz. Std	451.95	81.07	83.13	538.45
Minimum / Minimo	12.02	17.34	13.07	12.00
25th Percentile (P25)	12.03	62.03	62.95	12.02
Median / Mediana (P50)	12.03	135.34	136.31	12.03
75th Percentile (P75)	15.04	140.69	145.62	44.05
Maximum / Massimo	1292.00	268.70	272.44	1548.00
Variance / Varianza	204260.66	6571.68	6910.02	289928.78
90th Percentile (P90)	404.40	179.25	183.97	562.40


```
/var/folders/qn/y4fdcgw97jsbxdyb6s0d79680000gn/T/ipykernel_3998/1159961139.p
y:59: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(
/var/folders/qn/y4fdcgw97jsbxdyb6s0d79680000gn/T/ipykernel_3998/1159961139.p
y:80: UserWarning: set_ticklabels() should only be used with a fixed number
of ticks, i.e. after set_ticks() or using a FixedLocator.
```

```
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=15, ha='right')
/var/folders/qn/y4fdcgw97jsbxdyb6s0d79680000gn/T/ipykernel_3998/1159961139.p
y:94: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(
/var/folders/qn/y4fdcgw97jsbxdyb6s0d79680000gn/T/ipykernel_3998/1159961139.p
y:116: UserWarning: set_ticklabels() should only be used with a fixed number
of ticks, i.e. after set_ticks() or using a FixedLocator.
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=15, ha='right')
```

